

# CS4248 Assignment 2

U096883L Shawn Tan

## 1 Introduction

This assignment requires us to build a Part-of-Speech (POS) tagger, using training data from part of the Penn Treebank. The method used in our approach has to employ a Hidden Markov Model (HMM). This entails learning from the training data a set of parameters required for the HMM. Table 1 shows the various sets of data we have to collect for this particular assignment.

In the following sections, we will explain in detail how individual aspects of the HMM was created, outlining some of the technical difficulties faced. We also experiment with two simple smoothing techniques, Laplace (add-one) and Witten-Bell smoothing. The two techniques will be evaluated according to their precision, recall and F1 measures.

## 2 Learning from `sents.train`

The `sents.train` dataset contains 39,832 lines, each word annotated with POS tags. In order to extract the relevant information we count the transitions.

In order to obtain  $p(s_i | s_{i-1})$ , we need to count the different number of times one tag is followed by another. For each line, we extract the POS for each token, leaving a list of POS

Name	Description
$V$	all unique words
$S$	all unique POS tags
$p(s_i   s_{i-1})$	transition probability from one POS tag to another
$p(w   s)$	probability of seeing a word given a POS tag

Table 1: Parameters for a POS tagger HMM

tags. We then prepend a ‘^’ at the beginning, and a ‘\$’ character at the end. This ensures that the probabilities for a given POS starting a sentence are also taken into account.

Another preprocessing step performed was to change all instances of numeric tokens, and replacing all digits with #. For example, ‘\$25.50’ will be converted to ‘\$##.##’.

Going through the file, we maintain a dictionary of the following items:

Name	Description
$C(q_{t-1}, q_t)$	The number of times a POS tag occurs at time $t$ after a POS tag occurs at time $t - 1$
$C(q_t, q_{t-1})$	Count of transitions in reverse.
$C(q, w)$	Count of words given a POS tag.
$C(w, q)$	Count of POS tags given a word.

With this data, we can now perform different smoothing methods. The smoothing method is modular in our implementation, so as long as the counts are present, we can calculate the smoothed probabilities.

## 3 Evaluation

We performed evaluation for the POS tagger using two types of smoothing methods, Laplace smoothing using  $B = 1$  (or add-one smoothing) and Witten-Bell smoothing. The tagger performs smoothing for both conditional probabilities,  $p(w | q)$  and  $p(q_t | q_{t-1})$ .

### 3.1 Add-one smoothing

Using the add-one smoothing, we run 10-fold cross validation using `sents.out`. The results are shown in Table 2.

The results show that add-one smoothing causes a bad recall rate. This means that some of the words are tagged erroneously, but generally distributed throughout the other tags, such that they do not affect the other recall measures much. This suggests that the probability distributions given after smoothing spread the density out over the words and part of speech tags too much.

### 3.2 Witten-Bell smoothing

The results are shown in Table 3.

	Recall	Precision	$F1$
Fold 1	0.8645	0.7569	0.8618
Fold 2	0.8722	0.7702	0.8775
Fold 3	0.8726	0.7776	0.8656
Fold 4	0.8794	0.7599	0.8322
Fold 5	0.8924	0.7920	0.8626
Fold 6	0.8994	0.7603	0.8477
Fold 7	0.9110	0.7822	0.8523
Fold 8	0.8408	0.7705	0.8717
Fold 9	0.8689	0.7698	0.8462
Fold 10	0.8608	0.7811	0.8696
Ave.:			
Weighted ave.:			

Table 2: 10-fold validation using add-one smoothing

	Recall	Precision	$F1$
Fold 1	0.8541	0.8414	0.8719
Fold 2	0.8715	0.8978	0.8697
Fold 3	0.8595	0.9022	0.8723
Fold 4	0.8645	0.8719	0.8837
Fold 5	0.8890	0.8899	0.9024
Fold 6	0.8890	0.8766	0.8802
Fold 7	0.8924	0.8944	0.8877
Fold 8	0.8479	0.8773	0.8520
Fold 9	0.8468	0.8780	0.8736
Fold 10	0.8677	0.8929	0.8769
Ave.:			
Weighted ave.:			

Table 3: 10-fold validation using Witten-Bell smoothing